

# 3D Vision and Machine Perception

Prof. Kyungdon Joo

3D Vision & Robotics Lab.

AI Graduate School (AIGS) & Computer Science and Engineering (CSE)

# Contents

- Other types of cameras
- Image Processing Basic
- Image Gradients

# Other types of cameras

Imaging sensors

# 360 cameras

- capture all of the surrounding area around the camera without blind spots



360° video example

Samsung Gear 360




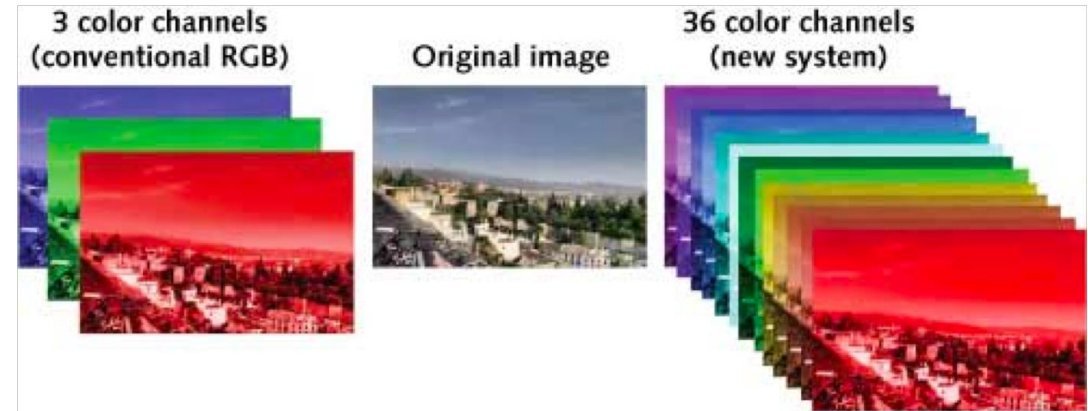
LG 360 Cam



Portable 360° cameras


# Multi-spectral cameras

- captures image data within specific wavelength ranges across the electromagnetic spectrum.




# Event camera

- Asynchronous events (pixel-level relative brightness changes caused by movement)
  - Low latency
  - Micro-second temporal resolution
  - High dynamic range (low-light to very bright)



Limits of standard camera




# Event camera-based vision applications

- Image Reconstruction from events



**Event Camera & Scene**




**Visualisation of Events**



# Event camera-based vision applications

- 6DoF Tracking from Photometric Map


Event camera



Standard camera



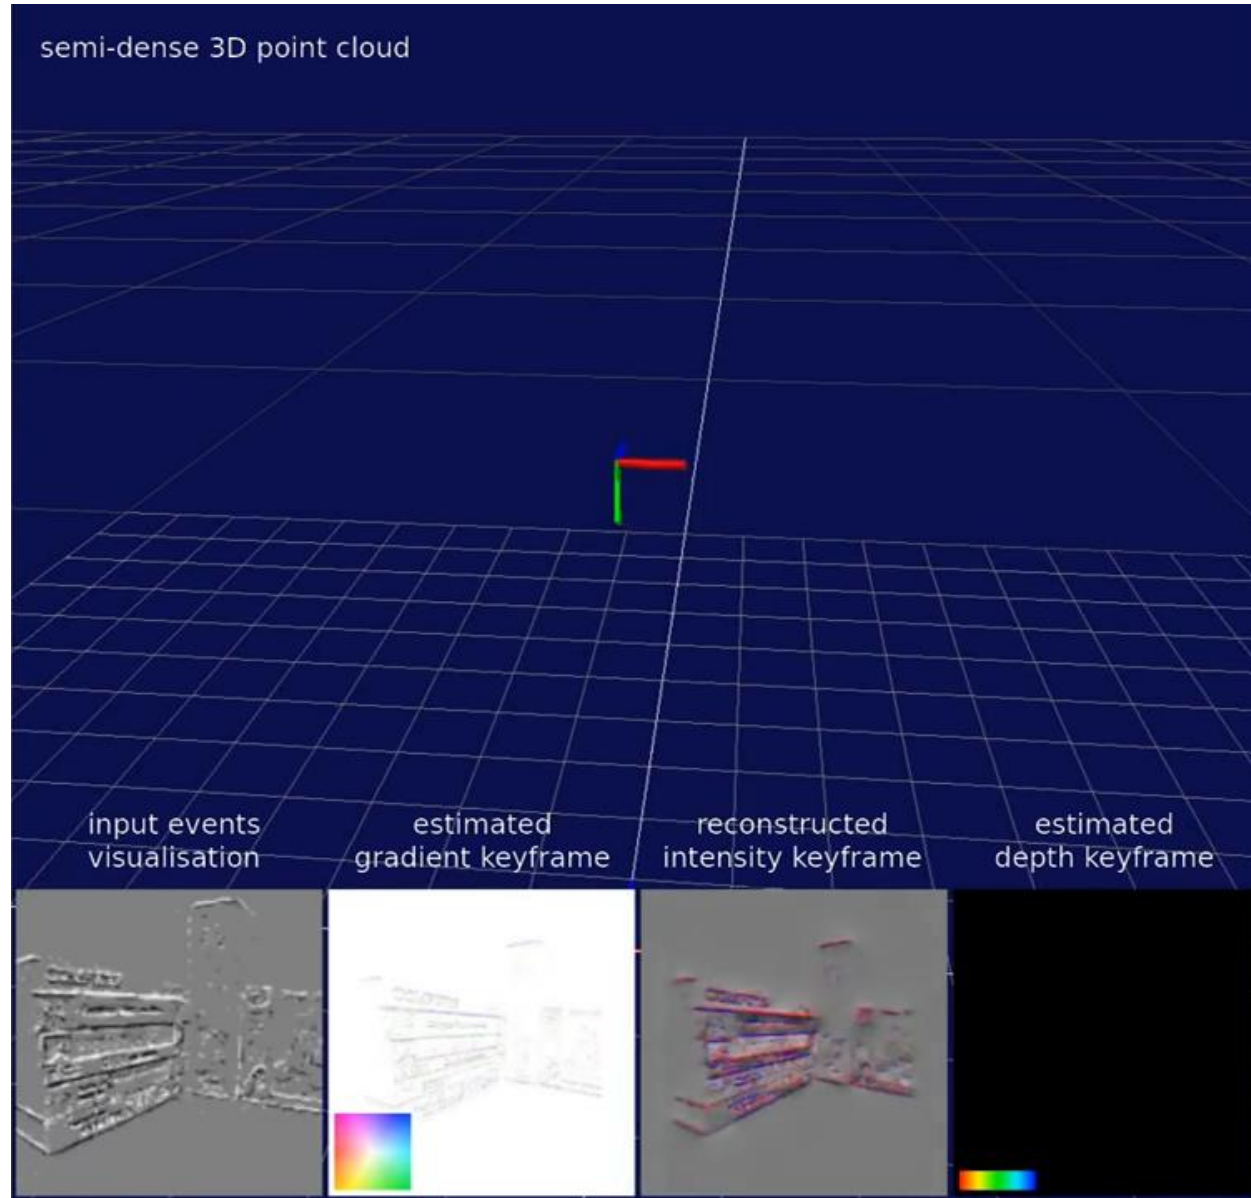
Motion estimation






# Event camera-based vision applications

- SLAM



# Event camera-based vision applications

- Dynamic Obstacle Detection & Avoidance
  - Works with relative speeds of up to 10 m/s
  - Perception latency: 3.5 ms



Falanga et al., Dynamic Obstacle Avoidance for Quadrotors with Event Cameras, Science Robotics, 2020.

Falanga et al. How Fast is too fast? The role of perception latency in high speed sense and avoid, RAL'19.

# Image Processing Basic

# What types of image transformations can we do?




Filtering



changes pixel *values*



Warping



changes pixel *locations*

# What types of image transformations can we do?

$F$



Filtering



$$G(\mathbf{x}) = h\{F(\mathbf{x})\}$$

$G$



changes *range* of image function

$F$




Warping



$$G(\mathbf{x}) = F(h\{\mathbf{x}\})$$


$G$



changes *domain* of image function


# What types of image filtering can we do?

Point operation



point processing

Neighborhood operation



“filtering”



Point processing

# Examples of point processing

original



darken



lower contrast



non-linear lower contrast



invert



lighten



raise contrast



non-linear raise contrast



# Examples of point processing

- How would you implement these?

original



darken



lower contrast



non-linear lower contrast



invert



lighten



raise contrast



non-linear raise contrast



# Examples of point processing

- How would you implement these?

original



$$x$$

darken



$$x - 128$$

lower contrast



$$\frac{x}{2}$$

invert



$$255 - x$$

lighten



$$x + 128$$

raise contrast



$$x \times 2$$

non-linear lower contrast



$$\left(\frac{x}{255}\right)^{1/3} \times 255$$

non-linear raise contrast



$$\left(\frac{x}{255}\right)^2 \times 255$$

Linear shift-invariant image filtering

# Linear shift-invariant image filtering





- Replace each pixel by a *linear* combination of its neighbors (and possibly itself).
- The combination is determined by the filter's *kernel*.
- The same kernel is *shifted* to all pixel locations so that all pixels use the same linear combination of their neighbors => *shift-invariant*.



# Convolution for 2D discrete signals

- Definition of filtering as convolution:

$$(f * g)(x, y) = \sum_{i, j = -\infty}^{\infty} f(i, j) I(x - i, y - j)$$


filtered image   filter  input image  notice the flip

# Convolution vs correlation

- Definition of discrete 2D convolution:

$$(f * g)(x, y) = \sum_{i, j=-\infty}^{\infty} f(i, j) I(x - i, y - j)$$


notice the flip



- Definition of discrete 2D correlation:

$$(f * g)(x, y) = \sum_{i, j=-\infty}^{\infty} f(i, j) I(x + i, y + j)$$

notice the lack of a flip




- Most of the time won't matter, because our kernels will be symmetric.

# The Gaussian filter

- named (like many other things) after Carl Friedrich Gauss
- kernel values sampled from the 2D Gaussian function:

$$f(i, j) = \frac{1}{2\pi\sigma^2} e^{-\frac{i^2+j^2}{2\sigma^2}}$$


- weight falls off with distance from center pixel
- theoretically infinite, in practice truncated to some maximum distance
- Any heuristics for selecting where to truncate?



# The Gaussian filter

- named (like many other things) after Carl Friedrich Gauss
- kernel values sampled from the 2D Gaussian function:

$$f(i, j) = \frac{1}{2\pi\sigma^2} e^{-\frac{i^2+j^2}{2\sigma^2}}$$




- weight falls off with distance from center pixel
- theoretically infinite, in practice truncated to some maximum distance
- Any heuristics for selecting where to truncate?
  - usually at  $2\sigma$  or  $3\sigma$

kernel  $\frac{1}{16}$


1	2	1
2	4	2
1	2	1

# Gaussian filtering example




# Gaussian vs box filtering

- Which blur do you think better?



original



7x7 Gaussian



7x7 box



# Other filters

input



filter

0	0	0
0	1	0
0	0	0

output



unchanged

# Other filters

input



filter

0	0	0
0	1	0
0	0	0

output



unchanged

input



filter

0	0	0
0	0	1
0	0	0

output

?

# Other filters


input



filter

0	0	0
0	1	0
0	0	0

output



unchanged

input



filter

0	0	0
0	0	1
0	0	0

output



shift to left  
by one

# Other filters

input



filter

0	0	0
0	2	0
0	0	0

$-\frac{1}{9}$


1	1	1
1	1	1
1	1	1

output



?

# Other filters

- do nothing for flat areas
- stress intensity peaks



# Sharpening examples





# Sharpening examples

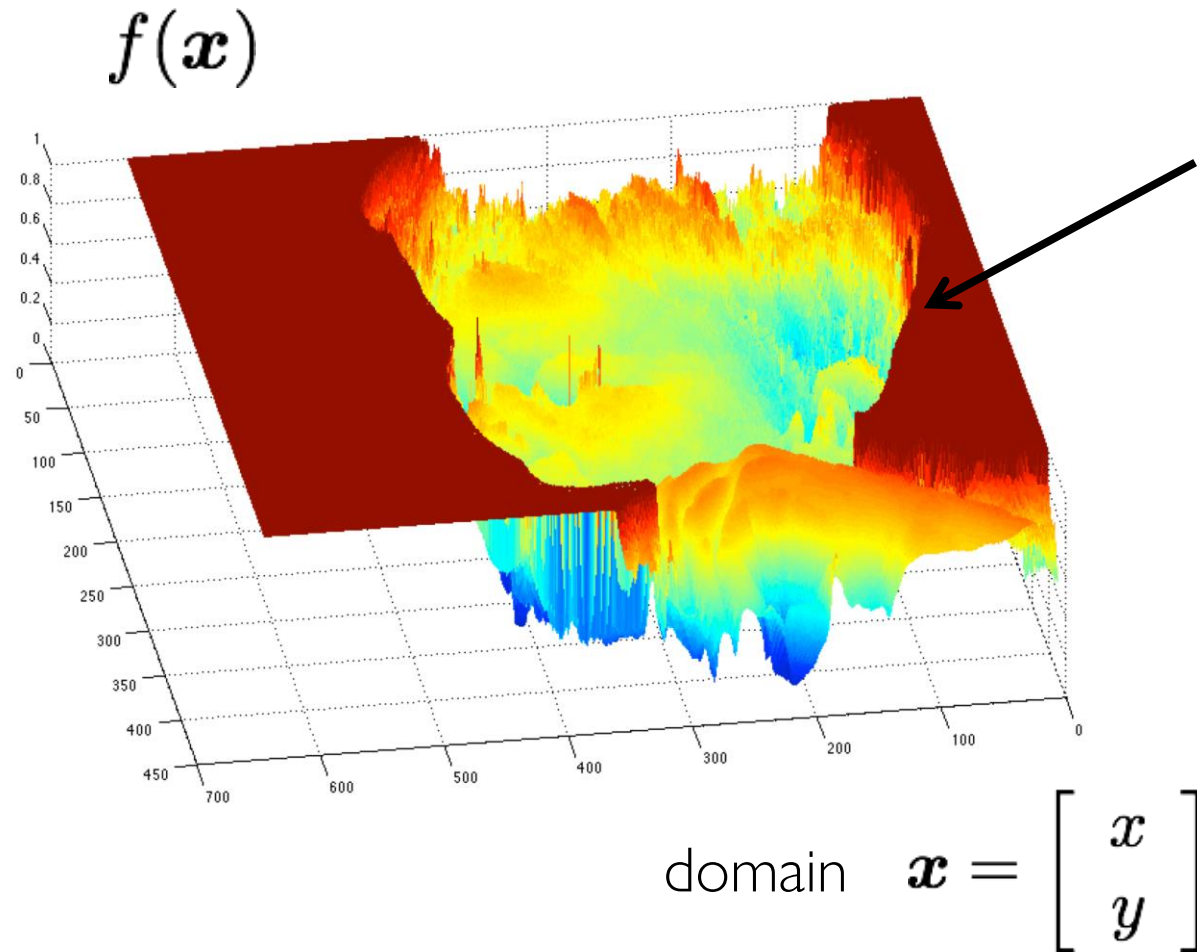


# Image Gradients

# What are image edges?



grayscale image



Very sharp  
discontinuities in  
intensity.

# Detecting edges

- How would you detect edges in an image  
(i.e., discontinuities in a function)?

# Detecting edges

- How would you detect edges in an image  
(i.e., discontinuities in a function)?
  - You take derivatives: derivatives are large at discontinuities.
- How do you differentiate a discrete image (or any other discrete signal)?

# Detecting edges

- How would you detect edges in an image  
(i.e., discontinuities in a function)?
  - You take derivatives: derivatives are large at discontinuities.
- How do you differentiate a discrete image (or any other discrete signal)?
  - You use finite differences.



# Finite differences

- High-school reminder: definition of a derivative using forward difference

$$f'(x) = \lim_{h \rightarrow 0} \frac{f(x+h) - f(x)}{h}$$

# Finite differences

- High-school reminder: definition of a derivative using forward difference

$$f'(x) = \lim_{h \rightarrow 0} \frac{f(x+h) - f(x)}{h}$$

- Alternative: use central difference

$$f'(x) = \lim_{h \rightarrow 0} \frac{f(x+0.5h) - f(x-0.5h)}{h}$$

- For discrete signals: Remove limit and set  $h = 2$

$$f'(x) = \frac{f(x+1) - f(x-1)}{2}$$

- What convolution kernel does this correspond to?

# Finite differences

- High-school reminder: definition of a derivative using forward difference

$$f'(x) = \lim_{h \rightarrow 0} \frac{f(x+h) - f(x)}{h}$$

- Alternative: use central difference

$$f'(x) = \lim_{h \rightarrow 0} \frac{f(x+0.5h) - f(x-0.5h)}{h}$$

- For discrete signals: Remove limit and set  $h = 2$

$$f'(x) = \frac{f(x+1) - f(x-1)}{2}$$

- What convolution kernel does this correspond to?

-1	0	1	?
----	---	---	---

1	0	-1	?
---	---	----	---

# Finite differences

- High-school reminder: definition of a derivative using forward difference

$$f'(x) = \lim_{h \rightarrow 0} \frac{f(x+h) - f(x)}{h}$$

- Alternative: use central difference

$$f'(x) = \lim_{h \rightarrow 0} \frac{f(x+0.5h) - f(x-0.5h)}{h}$$

- For discrete signals: Remove limit and set  $h = 2$

$$f'(x) = \frac{f(x+1) - f(x-1)}{2}$$

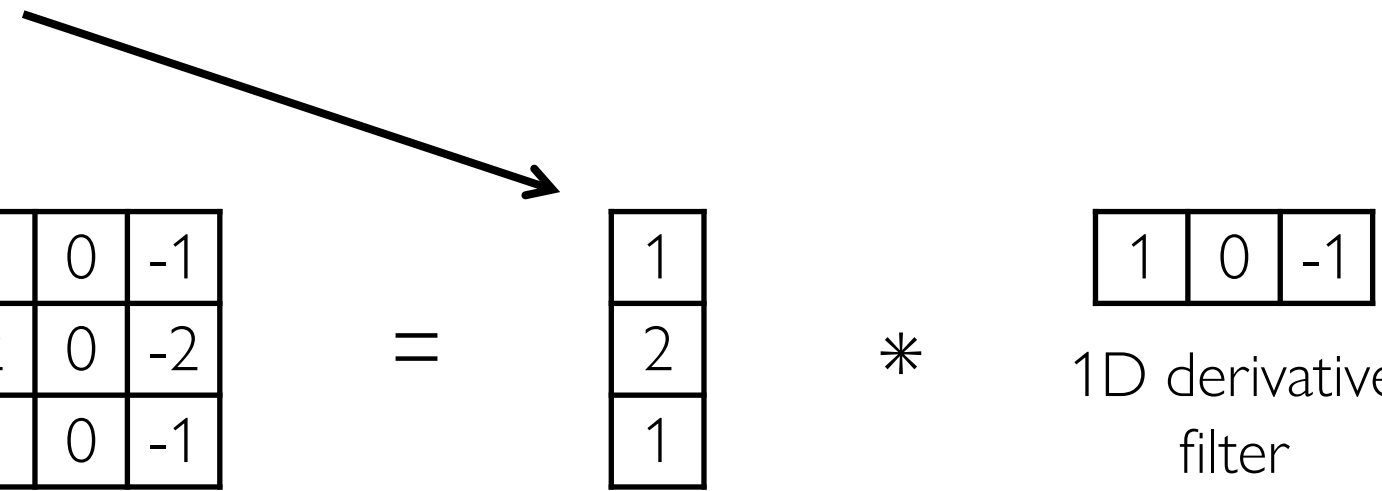
- What convolution kernel does this correspond to?

1D derivative filter

1	0	-1
---	---	----

# The Sobel filter

- What filter is this?



1	0	-1
2	0	-2
1	0	-1

Sobel filter

=

1
2
1

\*

1	0	-1
---	---	----

1D derivative  
filter

# The Sobel filter

- In a 2D image, does this filter responses along horizontal or vertical lines?

1	0	-1
2	0	-2
1	0	-1

Sobel filter

=

1
2
1

Blurring

\*

1	0	-1
---	---	----

1D derivative  
filter

# The Sobel filter

- Horizontal Sobel filter:

$$\begin{bmatrix} 1 & 0 & -1 \\ 2 & 0 & -2 \\ 1 & 0 & -1 \end{bmatrix} = \begin{bmatrix} 1 \\ 2 \\ 1 \end{bmatrix} * \begin{bmatrix} 1 & 0 & -1 \end{bmatrix}$$

- What does the vertical Sobel filter look like?



# The Sobel filter


- Horizontal Sobel filter (x-axis):

$$\begin{bmatrix} 1 & 0 & -1 \\ 2 & 0 & -2 \\ 1 & 0 & -1 \end{bmatrix} = \begin{bmatrix} 1 \\ 2 \\ 1 \end{bmatrix} * \begin{bmatrix} 1 & 0 & -1 \end{bmatrix}$$

- Vertical Sobel filter (y-axis):

$$\begin{bmatrix} 1 & 2 & 1 \\ 0 & 0 & 0 \\ -1 & -2 & -1 \end{bmatrix} = \begin{bmatrix} 1 \\ 0 \\ -1 \end{bmatrix} * \begin{bmatrix} 1 & 2 & 1 \end{bmatrix}$$

# Sobel filter example



original



horizontal Sobel filter



vertical Sobel filter

# Computing image gradients

- Select your favorite derivative filters.

$$\mathbf{S}_x = \begin{array}{|c|c|c|} \hline 1 & 0 & -1 \\ \hline 2 & 0 & -2 \\ \hline 1 & 0 & -1 \\ \hline \end{array}$$

$$\mathbf{S}_y = \begin{array}{|c|c|c|} \hline 1 & 2 & 1 \\ \hline 0 & 0 & 0 \\ \hline -1 & -2 & -1 \\ \hline \end{array}$$

# Computing image gradients

- Select your favorite derivative filters.

$$\mathbf{S}_x = \begin{array}{|c|c|c|} \hline 1 & 0 & -1 \\ \hline 2 & 0 & -2 \\ \hline 1 & 0 & -1 \\ \hline \end{array}$$

$$\mathbf{S}_y = \begin{array}{|c|c|c|} \hline 1 & 2 & 1 \\ \hline 0 & 0 & 0 \\ \hline -1 & -2 & -1 \\ \hline \end{array}$$

- Convolve with the image to compute derivatives.

$$\frac{\partial \mathbf{f}}{\partial x} = \mathbf{S}_x \otimes \mathbf{f}$$

$$\frac{\partial \mathbf{f}}{\partial y} = \mathbf{S}_y \otimes \mathbf{f}$$

# Computing image gradients

- Select your favorite derivative filters.

$$\mathbf{S}_x = \begin{array}{|c|c|c|} \hline 1 & 0 & -1 \\ \hline 2 & 0 & -2 \\ \hline 1 & 0 & -1 \\ \hline \end{array}$$

$$\mathbf{S}_y = \begin{array}{|c|c|c|} \hline 1 & 2 & 1 \\ \hline 0 & 0 & 0 \\ \hline -1 & -2 & -1 \\ \hline \end{array}$$

- Convolve with the image to compute derivatives.

$$\frac{\partial \mathbf{f}}{\partial x} = \mathbf{S}_x \otimes \mathbf{f}$$

$$\frac{\partial \mathbf{f}}{\partial y} = \mathbf{S}_y \otimes \mathbf{f}$$

- Form the image gradient and compute its direction and amplitude.

$$\nabla \mathbf{f} = \left[ \frac{\partial \mathbf{f}}{\partial x}, \frac{\partial \mathbf{f}}{\partial y} \right]$$

gradient

$$\theta = \tan^{-1} \left( \frac{\partial \mathbf{f}}{\partial y} / \frac{\partial \mathbf{f}}{\partial x} \right)$$

direction

$$\|\nabla \mathbf{f}\| = \sqrt{\left( \frac{\partial \mathbf{f}}{\partial x} \right)^2 + \left( \frac{\partial \mathbf{f}}{\partial y} \right)^2}$$


amplitude



# Image gradient example

- How does the gradient direction relate to these edges?


original



vertical  
derivative



gradient  
amplitude




horizontal  
derivative



# How do you find the edge of this signal?


intensity plot






# How do you find the edge of this signal?

intensity plot



- Using a derivative filter:

derivative plot




- What's the problem here?


# Differentiation is very sensitive to noise

- When using derivative filters, it is critical to blur first!


input




Gaussian



blurred




derivative of  
blurred



# Derivative of Gaussian (DoG) filter

- Derivative theorem of convolution:  $\frac{\partial}{\partial x}(h \star f) = (\frac{\partial}{\partial x}h) \star f$
- How many operations did we save?

input



derivative of  
Gaussian

output (same as  
before)

# Laplace filter

- Basically a second derivative filter.
  - We can use finite differences to derive it, as with first derivative filter.

first-order  
finite difference

$$f'(x) = \lim_{h \rightarrow 0} \frac{f(x + 0.5h) - f(x - 0.5h)}{h}$$



1D derivative filter

1	0	-1
---	---	----

second-order  
finite difference

$$f''(x) = \lim_{h \rightarrow 0} \frac{f(x + h) - 2f(x) + f(x - h)}{h^2}$$



Laplace filter

?

# Laplace filter

- Basically a second derivative filter.
  - We can use finite differences to derive it, as with first derivative filter.

first-order  
finite difference

$$f'(x) = \lim_{h \rightarrow 0} \frac{f(x + 0.5h) - f(x - 0.5h)}{h}$$



1D derivative filter

1	0	-1
---	---	----

second-order  
finite difference

$$f''(x) = \lim_{h \rightarrow 0} \frac{f(x + h) - 2f(x) + f(x - h)}{h^2}$$




Laplace filter

1	-2	1
---	----	---


# Laplacian of Gaussian (LoG) filter

- As with derivative, we can combine Laplace filtering with Gaussian filtering


input



Laplacian of  
Gaussian




output




# Laplacian of Gaussian (LoG) filter

- As with derivative, we can combine Laplace filtering with Gaussian filtering


input



Laplacian of  
Gaussian



output



“zero crossings” at edges



# Laplace and LoG filtering examples



Laplacian of Gaussian filtering



Laplace filtering

# Laplacian of Gaussian vs Derivative of Gaussian




Laplacian of Gaussian filtering




Derivative of Gaussian filtering

# Laplacian of Gaussian vs Derivative of Gaussian

- Zero crossings are more accurate at localizing edges (but not very convenient).




Laplacian of Gaussian filtering




Derivative of Gaussian filtering

# 2D Gaussian filters




$$h_{\sigma}(u, v) = \frac{1}{2\pi\sigma^2} e^{-\frac{u^2+v^2}{2\sigma^2}}$$

Gaussian




$$\frac{\partial}{\partial x} h_{\sigma}(u, v)$$




Derivative of Gaussian



x-direction



y-direction



Laplacian of Gaussian